

Writing Your First OpenCL Application

Zhongliang Chen

3-6-2015

Agenda

- OpenCL models and objects (65 min.)
 - Lecture (35 min.)
 - Exercises (30 min.)
- Profiling and debugging (25 min.)
 - Lecture (15 min.)
 - Exercise (10 min.)

OpenCL Architecture

- Portable parallel computing programming model
 - CPUs, GPUs, FPGAs, DSPs, etc.
- 4 models
 - Platform model
 - Execution model
 - Memory model
 - Programming model

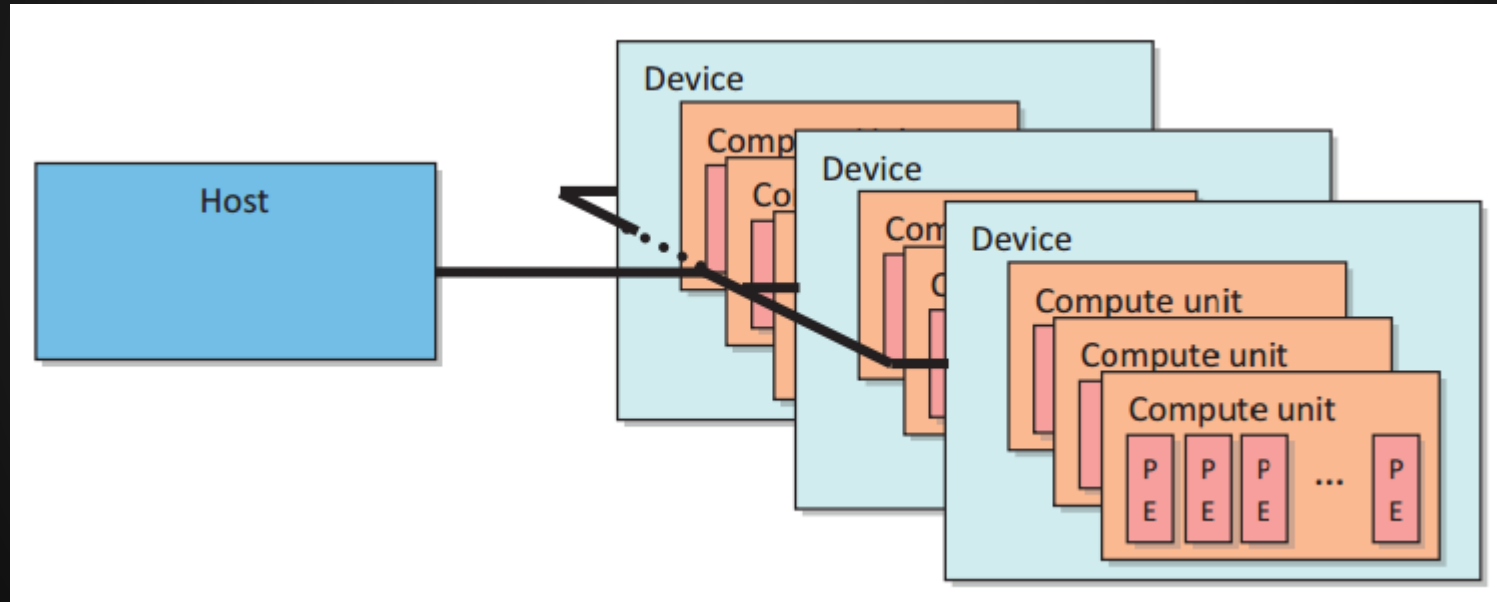
OpenCL Architecture

- 7 common objects
 - Platform
 - Context
 - Device
 - Command queue
 - Program
 - Kernel
 - Memory object

Platform Model

- Defined in OpenCL implementation, i.e., an OpenCL library
- Enables the host to interact with devices
- Currently each vendor provides only a single platform per implementation.

Platform Model



Platform Model

- Host
 - Processors an OpenCL library runs on
- Device
 - Processors an OpenCL library talks to
 - CPU: all cores are combined into a single device.
 - GPU: a GPU is a single device.

Get a Platform

```
cl_int clGetPlatformIDs(cl_uint num_entries,  
                        cl_platform_id *platforms,  
                        cl_uint *num_platforms)
```

- Usually called twice
 - First call to get the number of platforms
 - Then space is allocated for the platform objects.
 - Second call to get a platform

Exercise

Implement GetPlatform()

```
cl_int clGetPlatformIDs(cl_uint num_entries,  
                        cl_platform_id *platforms,  
                        cl_uint *num_platforms)
```

Context

- Environment for managing OpenCL objects and resources
- All OpenCL objects except platforms are managed by a context.
 - Devices
 - Programs
 - Kernels
 - Memory objects
 - Command queues

Create a Context

- Created with a type of devices

```
cl_context clCreateContextFromType (const cl_context_properties *properties,  
                                   cl_device_type device_type,  
                                   void (CL_CALLBACK *pfn_notify) (const char *errinfo,  
                                   const void *private_info,  
                                   size_t cb,  
                                   void *user_data),  
                                   void *user_data,  
                                   cl_int *errcode_ret)
```

Create a Context

- Created with one or more devices

```
cl_context clCreateContext(const cl_context_properties *properties,  
                          cl_uint num_devices,  
                          const cl_device_id *devices,  
                          (void CL_CALLBACK *pfn_notify) (  
                          const char *errinfo,  
                          const void *private_info, size_t cb,  
                          void *user_data  
                          ),  
                          void *user_data,  
                          cl_int *errcode_ret)
```

Exercise

Implement CreateContext()

```
cl_context clCreateContextFromType (const cl_context_properties *properties,  
                                   cl_device_type device_type,  
                                   void (CL_CALLBACK *pfn_notify) (const char *errinfo,  
                                   const void *private_info,  
                                   size_t cb,  
                                   void *user_data),  
                                   void *user_data,  
                                   cl_int *errcode_ret)
```

Get Devices

```
cl_int clGetDeviceIDs ( cl_platform_id platform ,  
                        cl_device_type device_type ,  
                        cl_uint num_entries ,  
                        cl_device_id *devices ,  
                        cl_uint *num_devices )
```

- Usually called twice
 - First call to get the number of devices
 - Then space is allocated for the device objects.
 - Second call to get one or more devices

Exercise

Implement GetDevice()

```
cl_int clGetDeviceIDs (cl_platform_id platform ,  
                      cl_device_type device_type ,  
                      cl_uint num_entries ,  
                      cl_device_id *devices ,  
                      cl_uint *num_devices )
```

Command Queues

- Connecting host and device
- Each device has its own command queue.
- Commands
 - Synchronous
 - **Asynchronous**
 - **In-order execution**
 - Out-of-order execution

Command Queues

- In-order execution
 - Each command is executed after the previous one has finished.
- Out-of-order execution
 - Commands are executed as soon as they are ready with no guarantees of their ordering
 - Events are commonly used for synchronization.

Command Queues (v1.2)

```
cl_command_queue clCreateCommandQueue(cl_context context,  
                                       cl_device_id device,  
                                       cl_command_queue_properties properties,  
                                       cl_int *errcode_ret)
```

Command-Queue Properties	Description
CL_QUEUE_OUT_OF_ORDER_EXEC_MODE_ENABLE	Determines whether the commands queued in the command-queue are executed in-order or out-of-order. If set, the commands in the command-queue are executed out-of-order. Otherwise, commands are executed in-order. See note below for more information.
CL_QUEUE_PROFILING_ENABLE	Enable or disable profiling of commands in the command-queue. If set, the profiling of commands is enabled. Otherwise profiling of commands is disabled. See clGetEventProfilingInfo for more information.

Exercise

Implement CreateCommandQueue()

```
cl_command_queue clCreateCommandQueue(cl_context context,  
                                       cl_device_id device,  
                                       cl_command_queue_properties properties,  
                                       cl_int *errcode_ret)
```

Programs

- A collection of OpenCL kernels
 - Source code or pre-compiled binary
 - Can contain constant data and auxiliary functions
- To create a program object
 - Read in a source code string or a binary

Programs

```
cl_program clCreateProgramWithSource ( cl_context context,  
                                       cl_uint count,  
                                       const char **strings,  
                                       const size_t *lengths,  
                                       cl_int *errcode_ret)
```

```
cl_program clCreateProgramWithBinary ( cl_context context,  
                                       cl_uint num_devices,  
                                       const cl_device_id *device_list,  
                                       const size_t *lengths,  
                                       const unsigned char **binaries,  
                                       cl_int *binary_status,  
                                       cl_int *errcode_ret)
```

Programs

- To build a program
 - Specify target devices.
 - Pass in compiler flags.
 - Check for compilation errors.

```
cl_int clBuildProgram (cl_program program,  
                      cl_uint num_devices,  
                      const cl_device_id *device_list,  
                      const char *options,  
                      void (CL_CALLBACK *pfn_notify)(cl_program program, void *user_data),  
                      void *user_data)
```

Programs

- If a program fails to compile, OpenCL requires programmers to explicitly ask for compiler output

```
cl_int clGetProgramBuildInfo (cl_program program,  
                             cl_device_id device,  
                             cl_program_build_info param_name,  
                             size_t param_value_size,  
                             void *param_value,  
                             size_t *param_value_size_ret)
```

Exercise

Implement CreateProgram()

```
cl_program clCreateProgramWithSource (cl_context context,  
                                     cl_uint count,  
                                     const char **strings,  
                                     const size_t *lengths,  
                                     cl_int *errcode_ret)
```

```
cl_int clBuildProgram (cl_program program,  
                      cl_uint num_devices,  
                      const cl_device_id *device_list,  
                      const char *options,  
                      void (CL_CALLBACK *pfn_notify)(cl_program program, void *user_data),  
                      void *user_data)
```


Kernels

- Functions declared in a program and running on a device
- Created from a compiled program

```
cl_kernel clCreateKernel (cl_program program,  
                          const char *kernel_name,  
                          cl_int *errcode_ret)
```

Kernels

- Arguments are explicitly associated with a kernel.
- Arguments can be memory objects or individual values.

```
cl_int clSetKernelArg ( cl_kernel kernel,  
                        cl_uint arg_index,  
                        size_t arg_size,  
                        const void *arg_value )
```

Exercise

Implement `CreateKernel()` and `SetKernelArg()`

```
cl_kernel clCreateKernel ( cl_program program,  
                           const char *kernel_name,  
                           cl_int *errcode_ret)
```

```
cl_int clSetKernelArg ( cl_kernel kernel,  
                        cl_uint arg_index,  
                        size_t arg_size,  
                        const void *arg_value )
```

Runtime Compilation

- OpenCL compiles programs and creates kernels at run time, which may incur high overhead. So each operation has to be performed only once.

What Are Left?

- Execute a kernel
 - Execution model
- Create memory objects and transfer data between host and device
 - Memory model

Execution Model

- Thread structure in SIMD
 - Each work-item works on one part of a problem.

Thread structure

Vector addition

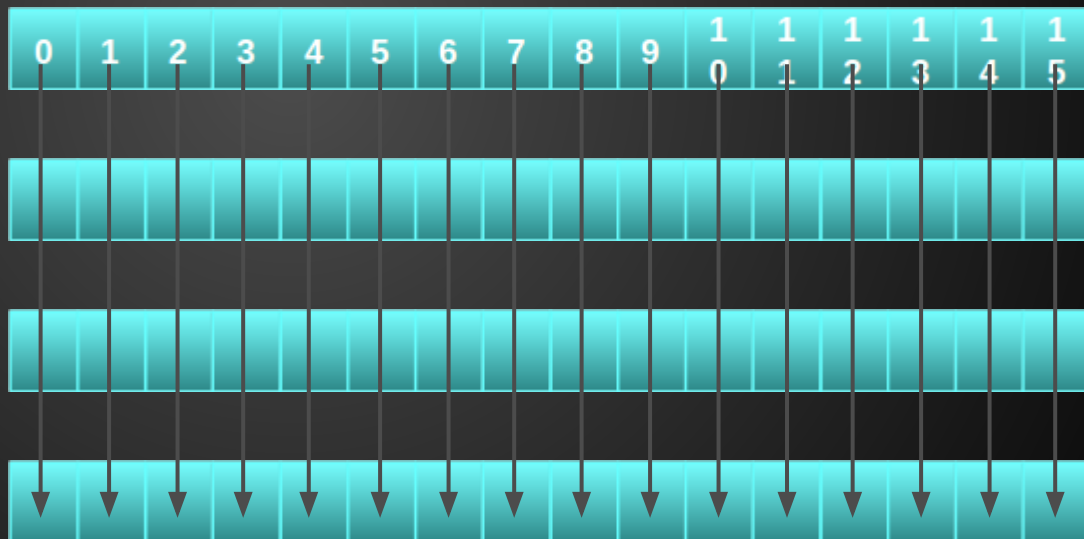
A

+

B

=

C



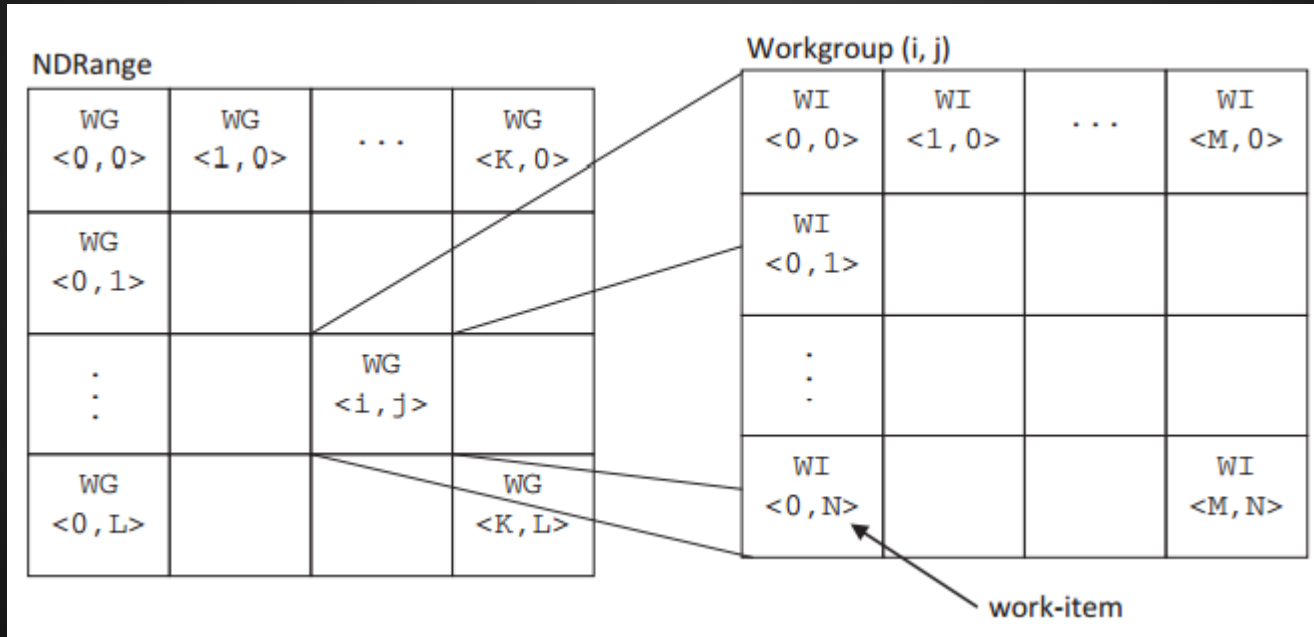
Execution Model

- Work-item: an instance of a kernel
- Work-group
 - Work-groups are independent of each other to guarantee scalability.
 - Scheduled to compute units and no guaranteed execution order
- Index space
 - A hierarchy of work-groups and work-items

Execution Model

- Work-items' IDs
 - Global ID, unique in the index space
 - Local ID in the work-group
 - $GID = WG_ID * WG_SIZE + LID$

Execution Model



Execution Model

<code>get_work_dim</code>	Number of dimensions in use
<code>get_global_size</code>	Number of global work items
<code>get_global_id</code>	Global work item ID value
<code>get_local_size</code>	Number of local work items
<code>get_local_id</code>	Local work item ID
<code>get_num_groups</code>	Number of work groups
<code>get_group_id</code>	Work group ID
<code>get_global_offset</code>	Work offset

Execution Model

```
cl_int clEnqueueNDRangeKernel ( cl_command_queue command_queue,  
                                cl_kernel kernel,  
                                cl_uint work_dim,  
                                const size_t *global_work_offset,  
                                const size_t *global_work_size,  
                                const size_t *local_work_size,  
                                cl_uint num_events_in_wait_list,  
                                const cl_event *event_wait_list,  
                                cl_event *event)
```

- Asynchronous execution
- A list of events can be used to specify prerequisite operations that must be completed before execution

```
cl_int clFinish ( cl_command_queue command_queue)
```

Exercise


Implement RunKernel()

```
cl_int clEnqueueNDRangeKernel ( cl_command_queue command_queue,  
                                cl_kernel kernel,  
                                cl_uint work_dim,  
                                const size_t *global_work_offset,  
                                const size_t *global_work_size,  
                                const size_t *local_work_size,  
                                cl_uint num_events_in_wait_list,  
                                const cl_event *event_wait_list,  
                                cl_event *event)
```


```
cl_int clFinish ( cl_command_queue command_queue)
```

Memory Model

- Global mem
- Constant mem
- Local mem
- Private mem



Memory Model



Memory Model

- Memory objects are explicitly managed.
 - Between host memory and device memory
 - Between global memory and local memory

```
cl_int clEnqueueWriteBuffer (cl_command_queue command_queue,  
                             cl_mem buffer,  
                             cl_bool blocking_write,  
                             size_t offset,  
                             size_t size,  
                             const void *ptr,  
                             cl_uint num_events_in_wait_list,  
                             const cl_event *event_wait_list,  
                             cl_event *event)
```

```
cl_int clEnqueueReadBuffer (cl_command_queue command_queue,  
                             cl_mem buffer,  
                             cl_bool blocking_read,  
                             size_t offset,  
                             size_t size,  
                             void *ptr,  
                             cl_uint num_events_in_wait_list,  
                             const cl_event *event_wait_list,  
                             cl_event *event)
```

Exercise

Implement WriteToGPU() and ReadFromGPU()

```
cl_int clEnqueueWriteBuffer (cl_command_queue command_queue,  
                             cl_mem buffer,  
                             cl_bool blocking_write,  
                             size_t offset,  
                             size_t size,  
                             const void *ptr,  
                             cl_uint num_events_in_wait_list,  
                             const cl_event *event_wait_list,  
                             cl_event *event)
```

```
cl_int clEnqueueReadBuffer (cl_command_queue command_queue,  
                             cl_mem buffer,  
                             cl_bool blocking_read,  
                             size_t offset,  
                             size_t size,  
                             void *ptr,  
                             cl_uint num_events_in_wait_list,  
                             const cl_event *event_wait_list,  
                             cl_event *event)
```


Writing a Kernel

- One instance of a kernel is created for each work-item.
- Kernels
 - Must begin with keyword `__kernel`
 - Must have return type `void`
 - Must declare the address space of each memory object argument
 - Use built-in work-item functions to get IDs, sizes, etc.

Exercise

Write VectorAddKernel

- Arguments: c, a, b, n
 - c : buffer, __global
 - a : buffer, __global
 - b : buffer, __global
 - n : unsigned int
- Functionality: vector addition $c = a + b$

Programming Model

- Data parallel
 - One-to-one mapping between work-items and elements in a memory object
 - Work-group defined explicitly or implicitly
- Task parallel
 - Kernel executed independent of an index space
- Synchronization
 - Between work-items in a work-group
 - Between command queues in a context

Wrap-Up

- 4 models
 - Platform, Execution, Memory, Programming
- 7 common objects
 - Platform, Context, Device
 - Command Queue, Program, Kernel, Memory Objects

Exercise

Make VectorAdd example work.

Agenda

- OpenCL models and objects (65 min.)
 - Lecture (35 min.)
 - Exercises (30 min.)
- Profiling and debugging (25 min.)
 - Lecture (15 min.)
 - Exercise (10 min.)

Command Queues With Profiling Enabled

```
cl_command_queue clCreateCommandQueue(cl_context context,  
                                       cl_device_id device,  
                                       cl_command_queue_properties properties,  
                                       cl_int *errcode_ret)
```

Command-Queue Properties	Description
CL_QUEUE_OUT_OF_ORDER_EXEC_MODE_ENABLE	Determines whether the commands queued in the command-queue are executed in-order or out-of-order. If set, the commands in the command-queue are executed out-of-order. Otherwise, commands are executed in order. See note below for more information.
CL_QUEUE_PROFILING_ENABLE	Enable or disable profiling of commands in the command-queue. If set, the profiling of commands is enabled. Otherwise profiling of commands is disabled. See clGetEventProfilingInfo for more information.

Events For Profiling

```
cl_int clEnqueueNDRangeKernel ( cl_command_queue command_queue,  
                                cl_kernel kernel,  
                                cl_uint work_dim,  
                                const size_t *global_work_offset,  
                                const size_t *global_work_size,  
                                const size_t *local_work_size,  
                                cl_uint num_events_in_wait_list,  
                                const cl_event *event_wait_list,  
                                cl_event *event)
```


Profiling API

```
cl_int clGetEventProfilingInfo (cl_event event,  
                                cl_profiling_info param_name,  
                                size_t param_value_size,  
                                void *param_value,  
                                size_t *param_value_size_ret)
```

cl_profiling_info	Return Type	Info. returned in <i>param_value</i>
CL_PROFILING_COMMAND_QUEUED	cl_ulong	A 64-bit value that describes the current device time counter in nanoseconds when the command identified by <i>event</i> is enqueued in a command-queue by the host.
CL_PROFILING_COMMAND_SUBMIT	cl_ulong	A 64-bit value that describes the current device time counter in nanoseconds when the command identified by <i>event</i> that has been enqueued is submitted by the host to the device associated with the command-queue.
CL_PROFILING_COMMAND_START	cl_ulong	A 64-bit value that describes the current device time counter in nanoseconds when the command identified by <i>event</i> starts execution on the device.
CL_PROFILING_COMMAND_END	cl_ulong	A 64-bit value that describes the current device time counter in nanoseconds when the command identified by <i>event</i> has finished execution on the device.

Exercise

Profile the kernel execution time

Demo

See how much a GPU is faster than a CPU.

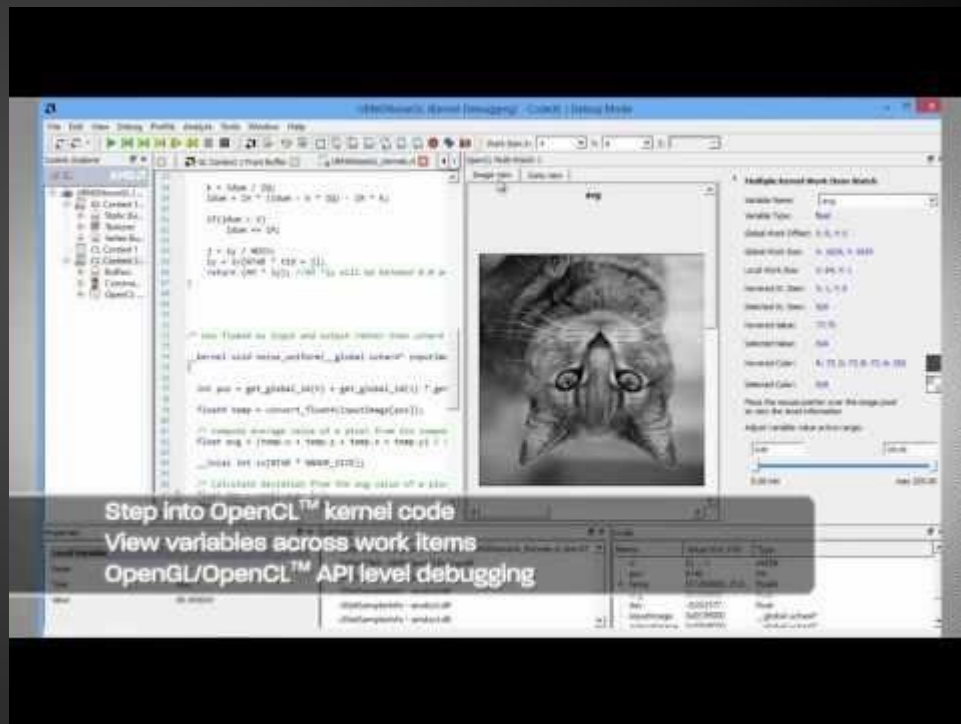
Debugging Kernels With `printf()`

- AMD GPUs support printing in kernels.
 - `cl_amd_printf` extension has to be enabled.
 - Remember that a kernel has many instances.

Generally one wants to check out only a few work-items.
- Information to be printed is buffered until the kernel completes and then transferred back to the host.
 - The kernel must be finished.

AMD CodeXL

- Integrated tool for debugging and profiling



Programming Multiple Devices

- Single context, multiple devices
 - Simple
 - Memory objects are common in a context
- Multiple context, multiple devices
 - Computing on a cluster
- Considerations for CPU-GPU heterogeneous computing
 - Scheduling, load balancing

Summary

- OpenCL models
 - Platform, Execution, Memory, Programming
- OpenCL common objects
 - Platform, Context, Device
 - Command Queue, Program, Kernel, Memory Object
- Profiling and debugging

Next Session (Leiming Yu)

- Global memory is off chip and so slow, which is a very bad idea for data reuse. Local memory is on chip and much faster. How to use it?
- Vector addition is just a toy. How to write and optimize a real application?

Reference

- OpenCL university kit
 - By Perhaad Mistry and Dana Schaa
- AMD OpenCL Programming User Guide
 - rev 1.0 Beta
- Heterogeneous Computing with OpenCL
 - By Benedict Gaster, Lee Howes, David R. Kaeli, Perhaad Mistry & Dana Schaa

Thank you!

Questions?

OpenCL 2.0 Main Features

- SVM: Shared Virtual Memory
 - Pointer-containing data structures can be easily shared between host and device.
- DP: Dynamic Parallelism
 - Kernels can enqueue kernels without host interactions.
- Pipes
 - Pipes enable data transfers between kernels without host interactions.