

# Matrix Multiplication using OpenCL

Leiming Yu

3-16-2015

# Previous Sections

- Introduction to Parallelism
  - CPU vs. GPU
- OpenCL compilation and framework
- OpenCL Case Study: vector add
  - Create objects, e.g., platform, context, program
  - Write kernel
  - Transfer data
  - Profile kernel execution

# Current Section

- Matrix Multiplication (MM) Background
- MM on GPU
  - Memory
  - Execution
  - Validation
  - Performance Benchmarking

# Matrix Multiplication

$$\begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{bmatrix} \times \begin{bmatrix} 7 & 8 \\ 9 & 10 \\ 11 & 12 \end{bmatrix} = \begin{bmatrix} 58 & \end{bmatrix}$$

The "Dot Product" is where we **multiply matching members**, then sum up:

$$(1, 2, 3) \cdot (7, 9, 11) = 1 \times 7 + 2 \times 9 + 3 \times 11 = 58$$

We match the 1st members (1 and 7), multiply them, likewise for the 2nd members (2 and 9) and the 3rd members (3 and 11), and finally sum them up.

# Matrix Multiplication

$$\begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{bmatrix} \times \begin{bmatrix} 7 & 8 \\ 9 & 10 \\ 11 & 12 \end{bmatrix} = \begin{bmatrix} 58 & 64 \\ 139 & 154 \end{bmatrix}$$

$$(1, 2, 3) \cdot (8, 10, 12) = 1 \times 8 + 2 \times 10 + 3 \times 12 = 64$$

$$\begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{bmatrix} \times \begin{bmatrix} 7 & 8 \\ 9 & 10 \\ 11 & 12 \end{bmatrix} = \begin{bmatrix} 58 & 64 \\ 139 & 154 \end{bmatrix} \quad \checkmark$$

DONE!

# An Example

Beef pies cost **\$3** each

Chicken pies cost **\$4** each

Vegetable pies cost **\$2** each

	Mon	Tue	Wed	Thu
Beef	13	9	7	15
Chicken	8	7	4	6
Vegetable	6	4	0	3

$$\begin{bmatrix} \$3 & \$4 & \$2 \end{bmatrix} \times \begin{bmatrix} 13 & 9 & 7 & 15 \\ 8 & 7 & 4 & 6 \\ 6 & 4 & 0 & 3 \end{bmatrix} = \begin{bmatrix} \$83 & \$63 & \$37 & \$75 \end{bmatrix}$$

$\$3 \times 13 + \$4 \times 8 + \$2 \times 6$

# More about MM

## Order of Multiplication


In arithmetic we are used to:

$$3 \times 5 = 5 \times 3$$

(The Commutative Law of Multiplication)

But this is **not** generally true for matrices (matrix multiplication is **not commutative**):

$$AB \neq BA$$




$$\begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix} \times \begin{bmatrix} 2 & 0 \\ 1 & 2 \end{bmatrix} = \begin{bmatrix} 4 & 4 \\ 10 & 8 \end{bmatrix}$$

$$\begin{bmatrix} 2 & 0 \\ 1 & 2 \end{bmatrix} \times \begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix} = \begin{bmatrix} 2 & 4 \\ 7 & 10 \end{bmatrix}$$

## Other interesting terms

- Identity Matrix
- Inverse of a Matrix
- Determinant of a Matrix

# MM on CPU



```
for (i=0; i<N; i++)  
  for (j=0; j<N; j++)  
    C[i, j] = 0;  
    for (k=0; k<N; k++)  
      C[i, j] += A[i, k]*B[k, j];
```



# Code Exercise

- CPU implementation of MM

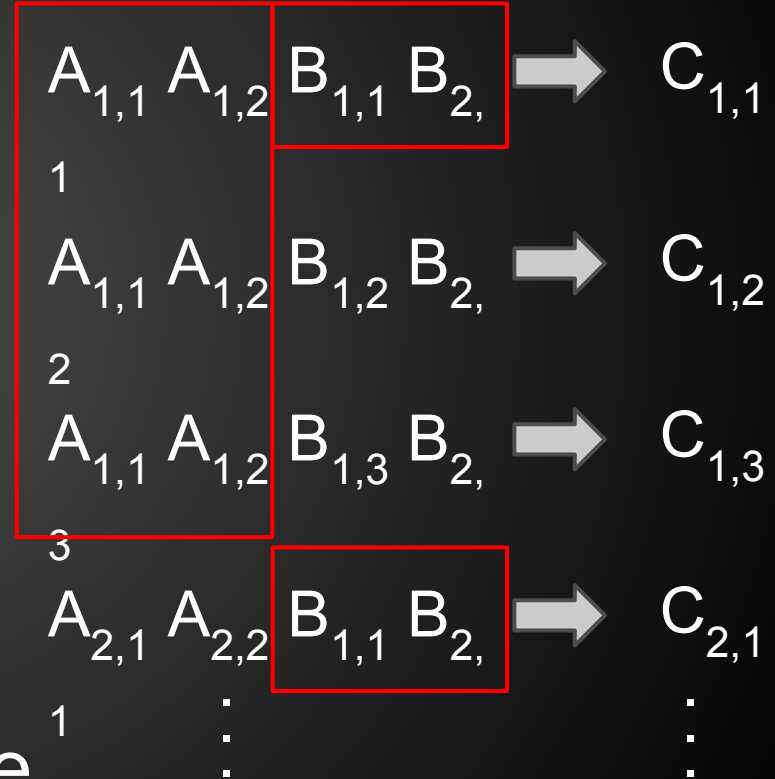

# Naive MM on GPU

- Launch 2D Grid
- Customize the 2D workgroup size
- Store data in global/device memory
- Each thread computes one element of the output matrix

# Code Exercise

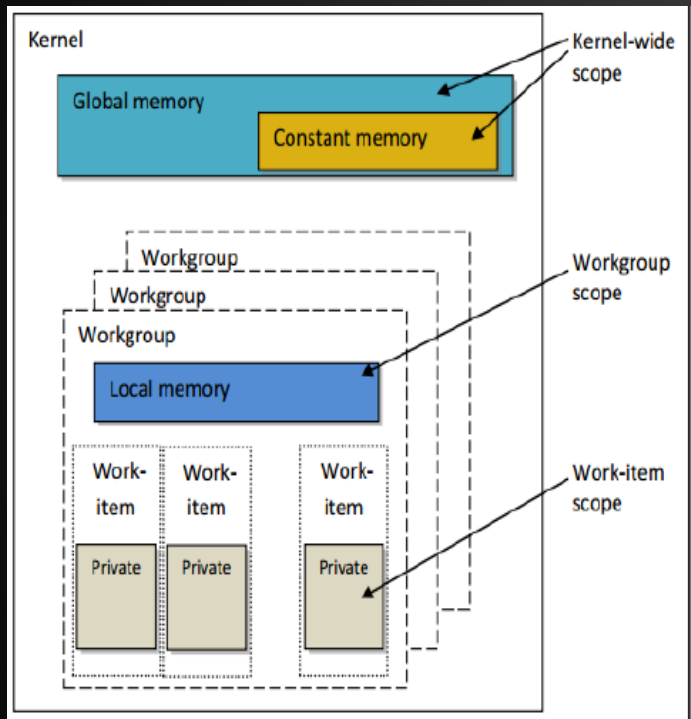
- CPU implementation of MM
- Naive MM implementation on GPU

# Problem with naive version




Data reuse

# Solution



- Local memory for data reuse
- Chunk the matrix into sub-matrices for parallel computation and good data locality

# Scheme



# Code Exercise

- CPU implementation of MM
- Naive MM implementation on GPU
- Optimized MM
  - Tiled computation
  - Local Memory

Congratulations!

You are leveled up ^~^

